

Artificial Intelligence

Lecture 3: Adversarial Search

Xiaojin Gong

2022-03-07


Credits: AI Courses in Berkeley

Review

- Problem-Solving Agents
 - Problem Formulation
 - Solving Problems by Searching
 - Uninformed Search
 - Breadth First Search
 - Depth First Search
 - Iterative Deepening Search
 - Cost-sensitive Search
 - Informed Search
 - Best First Search
 - A* Search


Outline

- Adversarial Search Problem (a.k.a Games)
 - Game Formulation
 - Minimax Search
 - Alpha-Beta Pruning
 - Monte Carlo Tree Search
- AlphaGo



Adversarial Search

- Adversarial search problems (a.k.a **Games**)
 - Multi-agent environment, competitive
 - Solving by adversarial search
- Types of Games
 - One, two, or more players?
 - Zero sum?
 - Perfect information (can you see the state)?
 - Deterministic or stochastic?




Deterministic Games

- A **search problem** consists of
 - S_0 : the initial state
 - $\text{Player}(s)$: defines which player has the move in a state
 - $\text{Actions}(s)$: returns the set of legal moves in a state
 - $\text{Result}(s, a)$: the transition model
 - $\text{Terminal-test}(s)$: is true when the game is over and false otherwise
 - $\text{Utility}(s, p)$: a utility / objective function
- A **solution** for a player is a **policy** which specifies each move.

Game Tree


- The initial state, actions, and results define the game tree




- $9! = 362880$ terminal nodes

Game Tree

- The initial state, actions, and results define the game tree




- $b \approx 35, m \approx 100$
- $O(b^m)$



- $b \approx 250, m \approx 150$
- $O(b^m)$


Minimax Search

- A game tree
- Players alternate turns
- Compute each node's **minimax** value:
 - The minimax value is **the utility** of being in the state, assuming that both players play optimally from there to the end of the game.



Minimax Search

$$\text{MINIMAX}(s) = \begin{cases} \text{UTILITY}(s) & \text{if } \text{TERMINAL-TEST}(s) \\ \max_{a \in \text{Actions}(s)} \text{MINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MAX} \\ \min_{a \in \text{Actions}(s)} \text{MINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MIN} \end{cases}$$




The Minimax Algorithm

```
function MINIMAX-DECISION(state) returns an action  
  return  $\arg \max_{a \in \text{ACTIONS}(s)} \text{MIN-VALUE}(\text{RESULT}(state, a))$ 
```

```
function MAX-VALUE(state) returns a utility value  
  if TERMINAL-TEST(state) then return UTILITY(state)  
   $v \leftarrow -\infty$   
  for each a in ACTIONS(state) do  
     $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a)))$   
  return v
```

```
function MIN-VALUE(state) returns a utility value  
  if TERMINAL-TEST(state) then return UTILITY(state)  
   $v \leftarrow \infty$   
  for each a in ACTIONS(state) do  
     $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a)))$   
  return v
```




Minimax Search

- Performance


- Performs a complete **depth-first** exploration
- Time complexity: $O(b^m)$
- Space complexity: $O(bm)$
- Complete? **Yes** (if tree is finite)
- Optimal? **Yes** (against an optimal opponent)

- Examples:


- **Chess**, $b \approx 35$, $m \approx 100$
- **Go game**, $b \approx 250$, $m \approx 150$



Alpha-Beta Pruning




Alpha-Beta Pruning



$$\begin{aligned}\text{MINIMAX}(\text{root}) &= \max(\min(3, 12, 8), \min(2, x, y), \min(14, 5, 2)) \\ &= \max(3, \min(2, x, y), 2) \\ &= \max(3, z, 2) \quad \text{where } z = \min(2, x, y) \leq 2 \\ &= 3.\end{aligned}$$

Alpha-Beta Pruning

- General principle:
 - Consider a **node n** ;
 - If MAX has a better choice α either at the parent node of n or at any choice point further up, then n will never be reached in actual play.
 - So once we found out enough about n (by **examining some of its descendants**) to reach this conclusion, we can prune it.
- Alpha-Beta Pruning
 - α = the highest value for MAX on path to root
 - β = the lowest value for MIN on path to root



Alpha-Beta Pruning

function ALPHA-BETA-SEARCH(*state*) **returns** an action
 $v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$
 return the *action* in ACTIONS(*state*) with value *v*

function MAX-VALUE(*state*, α , β) **returns** a utility value
 if TERMINAL-TEST(*state*) **then return** UTILITY(*state*)
 $v \leftarrow -\infty$
 for each *a* **in** ACTIONS(*state*) **do**
 $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$
 if $v \geq \beta$ **then return** *v*
 $\alpha \leftarrow \text{MAX}(\alpha, v)$
 return *v*


function MIN-VALUE(*state*, α , β) **returns** a utility value
 if TERMINAL-TEST(*state*) **then return** UTILITY(*state*)
 $v \leftarrow +\infty$
 for each *a* **in** ACTIONS(*state*) **do**
 $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$
 if $v \leq \alpha$ **then return** *v*
 $\beta \leftarrow \text{MIN}(\beta, v)$
 return *v*

Alpha-Beta Pruning







function ALPHA-BETA-SEARCH(*state*) **returns** an action
 $v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$
return the action in ACTIONS(*state*) with value *v*

function MAX-VALUE(*state*, α , β) **returns** a utility value
if TERMINAL-TEST(*state*) **then return** UTILITY(*state*)
 $v \leftarrow -\infty$
for each *a* **in** ACTIONS(*state*) **do**
 $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$
if $v \geq \beta$ **then return** *v*
 $\alpha \leftarrow \text{MAX}(\alpha, v)$
return *v*

function MIN-VALUE(*state*, α , β) **returns** a utility value
if TERMINAL-TEST(*state*) **then return** UTILITY(*state*)
 $v \leftarrow +\infty$
for each *a* **in** ACTIONS(*state*) **do**
 $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$
if $v \leq \alpha$ **then return** *v*
 $\beta \leftarrow \text{MIN}(\beta, v)$
return *v*



Alpha-Beta Pruning



Alpha-Beta Pruning

- The effectiveness of alpha-beta pruning is highly dependent on the order in which the states are examined.
- In perfect case, alpha-beta needs to examine only $O(b^{m/2})$ nodes, instead of $O(b^m)$ for minimax.
- In random order, alpha-beta needs to examine roughly $O(b^{3m/4})$ nodes.

Imperfect Real-time Decisions


- Problem: In realistic games, cannot search to leaves !
- Solution 1: Depth-limited search
 - Search only to a limited depth in the tree
 - Replace terminal utilities with a heuristic evaluation function for non-terminal positions

H-MINIMAX(s, d) =

$$\begin{cases} \text{EVAL}(s) & \text{if CUTOFF-TEST}(s, d) \\ \max_{a \in \text{Actions}(s)} \text{H-MINIMAX}(\text{RESULT}(s, a), d + 1) & \text{if PLAYER}(s) = \text{MAX} \\ \min_{a \in \text{Actions}(s)} \text{H-MINIMAX}(\text{RESULT}(s, a), d + 1) & \text{if PLAYER}(s) = \text{MIN.} \end{cases}$$

Evaluation Functions

- An **evaluation function** returns an estimate of the expected utility of the game from a given position.



- Ideal function: returns the actual minimax value of the position
- In practice: typically weighted linear sum of features:

$$\text{EVAL}(s) = w_1 f_1(s) + w_2 f_2(s) + \cdots + w_n f_n(s) = \sum_{i=1}^n w_i f_i(s)$$

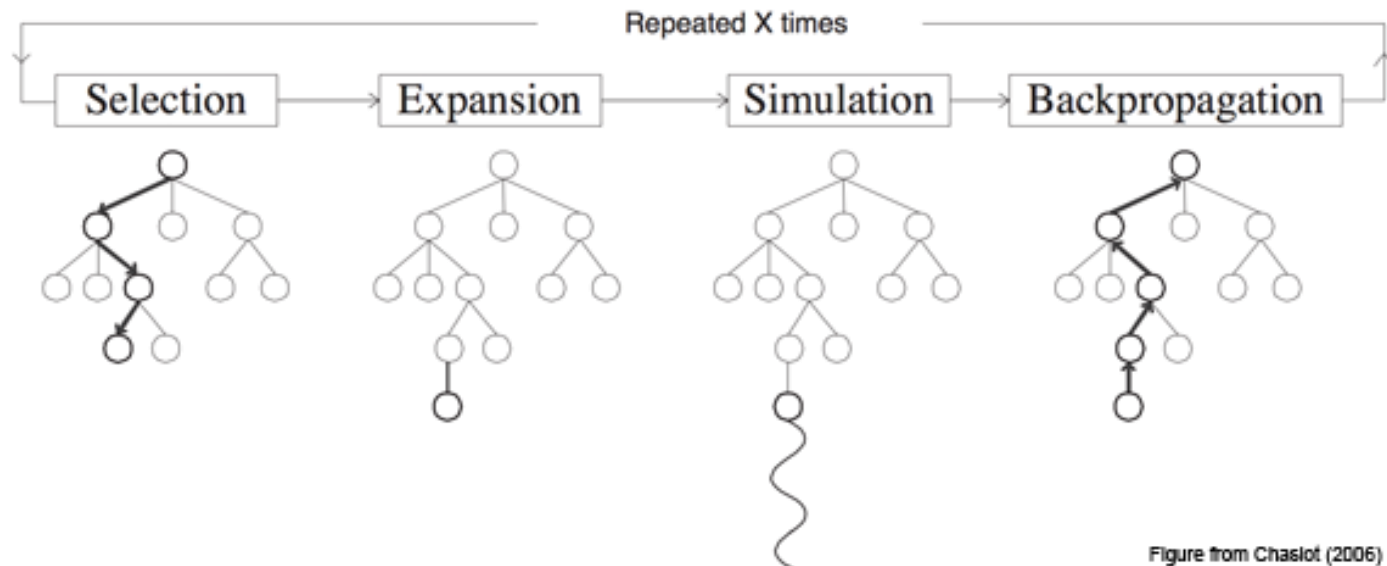
- E.g. $f_1(s) = (\text{num white queens} - \text{num black queens})$

Evaluation Functions

- Evaluation function are always imperfect
- The deeper in the tree the evaluation function is buried, the less the quality of the evaluation function matters
- An important example of the tradeoff between complexity of features and complexity of computation


Imperfect Real-time Decisions

- Problem: In realistic games, cannot search to leaves !
- Solution 2: Monte Carlo Tree Search
 - MTCS builds a statistics tree that partially maps onto the entire game tree
 - Statistics tree guides to “look only/mostly at the most interesting nodes in the game tree”
 - Value of nodes determined by simulations



Monte Carlo Tree Search

- Selection → Expansion → Simulation → Backpropagation




Selection

Starting at root node R, recursively select optimal child nodes until a leaf node L is reached.

Monte Carlo Tree Search

- Selection → Expansion → Simulation → Backpropagation

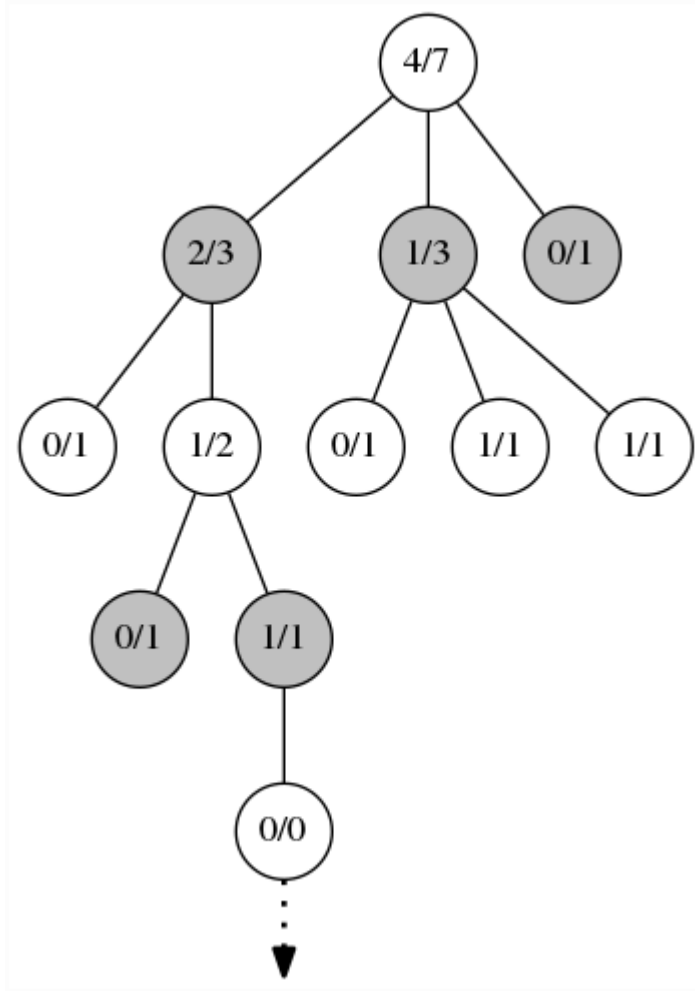


Expansion

If L is not a terminal node then create one or more child nodes and select one C.

Monte Carlo Tree Search

- Selection → Expansion → Simulation → Backpropagation





Simulation

Run a simulated playout from C until a result is achieved.

Rollout: Randomly choose an action at each step and simulate this action to receive an average reward when the game is over.

Monte Carlo Tree Search

- Selection → Expansion → Simulation → Backpropagation



Backpropagation

Update the current move sequence with the simulation result.


Adversarial Search

1. A. S. Douglas programmed the **first software that managed to master a game in 1952**. The game? Tic-Tac-Toe! This was part of his doctoral dissertation at Cambridge
2. *A few years later, Arthur Samuel was the first to use reinforcement learning that to play Checkers by playing against itself*
3. In 1992, Gerald Tesauro designed a now-popular **program called TD-Gammon to play backgammon at a world-class level**
4. For decades, Chess was seen as “the ultimate challenge of AI” . **IBM’ s Deep Blue was the first software that exhibited superhuman Chess capability**. The system famously defeated Garry Kasparov, the reigning grandmaster of chess, in 1997
5. **One of the most popular board game AI milestones was reached in 2016 in the game of Go**. Lee Sedol, a 9-dan professional Go player, lost a five-game match against Google DeepMind’ s AlphaGo software which featured a deep reinforcement learning approach
6. Notable recent milestones in video game AI include **algorithms developed by Google DeepMind to play several games from the classic Atari 2600** video game console at a super-human skill level
7. Last year, OpenAI built the popular [OpenAI Five](#) system that mastered the complex strategy game of DOTA


<https://www.analyticsvidhya.com/blog/2019/01/monte-carlo-tree-search-introduction-algorithm-deepmind-alphago/>

AlphaGo

2016. 3



2017. 5



ARTICLE

NATURE | VOL 529 | 28 JANUARY 2016

doi:10.1038/nature16961

Mastering the game of Go with deep neural networks and tree search

David Silver^{1*}, Aja Huang^{1*}, Chris J. Maddison¹, Arthur Guez¹, Laurent Sifre¹, George van den Driessche¹, Julian Schrittwieser¹, Ioannis Antonoglou¹, Veda Panneershelvam¹, Marc Lanctot¹, Sander Dieleman¹, Dominik Grewe¹, John Nham², Nal Kalchbrenner¹, Ilya Sutskever², Timothy Lillicrap¹, Madeleine Leach¹, Koray Kavukcuoglu¹, Thore Graepel¹ & Demis Hassabis¹

ARTICLE


NATURE | VOL 550 | 19 OCTOBER 2017

doi:10.1038/nature24270


Mastering the game of Go without human knowledge

David Silver^{1*}, Julian Schrittwieser^{1*}, Karen Simonyan^{1*}, Ioannis Antonoglou¹, Aja Huang¹, Arthur Guez¹, Thomas Hubert¹, Lucas Baker¹, Matthew Lai¹, Adrian Bolton¹, Yutian Chen¹, Timothy Lillicrap¹, Fan Hui¹, Laurent Sifre¹, George van den Driessche¹, Thore Graepel¹ & Demis Hassabis¹

AlphaGo



- $b \approx 35, m \approx 100$
- $O(b^m)$





- $b \approx 250, m \approx 150$
- $O(b^m)$

AlphaGo

- The core parts of AlphaGo:
 - Convolutional Neural Networks:
 - Evaluates new positions & moves
 - Reinforcement learning:
 - Trains the AI by using the current best agent to play against itself
 - Monte Carlo Tree Search:
 - Chooses the next move

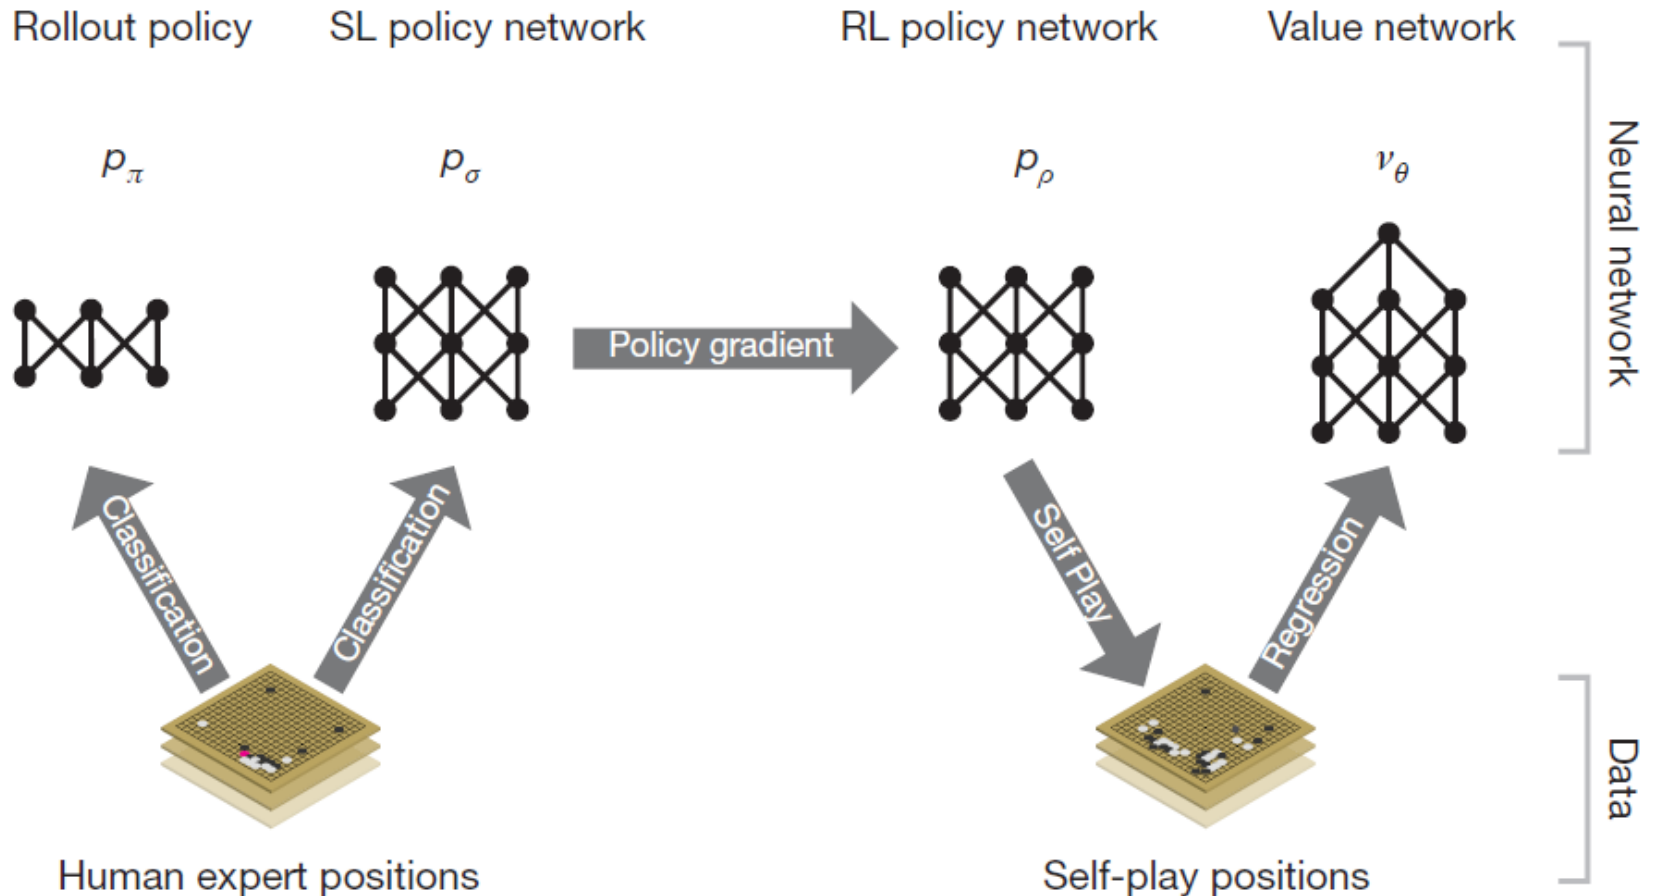
AlphaGo

- Two policies:
 - The depth of the search is reduced by position evaluation.
 - The breath of the search is reduced by sampling actions from a policy $p(a|s)$.
- Two networks:
 - Value network to evaluate board position.
 - Policy network to select moves.



AlphaGo

- Reinforcement learning



AlphaGo

- Monte Carlo tree search in AlphaGo





Figure 3 | Monte Carlo tree search in AlphaGo.

AlphaGo

- Monte Carlo tree search in AlphaGo

a

Selection




- Selecting the edge with maximum action value Q plus a bonus $u(P)$.

AlphaGo

- Monte Carlo tree search in AlphaGo


b Expansion



- The leaf node may be expanded; the new node is processed once by the policy network p_σ and the output probabilities are stored as prior probabilities P for each action.

AlphaGo

- Monte Carlo tree search in AlphaGo




- At the end of a simulation, the leaf node is evaluated in two ways:
 - using the value network v_θ ;
 - running a rollout to the end of the game with the fast rollout policy p_π ;
 - then computing the winner with function r .

AlphaGo

- Monte Carlo tree search in AlphaGo

d

Backup



- Action values Q are updated to track the mean value of all evaluations $r(\cdot)$ and $v_\theta(\cdot)$ in the subtree below that action.

$$Q(s, a) = \frac{1}{N(s, a)} \sum_{i=1}^n 1(s, a, i) V(s_L^i)$$

$$N(s, a) = \sum_{i=1}^n 1(s, a, i)$$

$$V(s_L) = (1 - \lambda)v_\theta(s_L) + \lambda z_L$$

$$u(s, a) \propto \frac{P(s, a)}{1 + N(s, a)}$$

$$P(s, a) = p_{\sigma}(a|s)$$

Assignment

- Reading assignment:
 - Ch. 5.1-5.5
- Homework 2:
 - Due by Mar. 14, 2022.
- Project 1:
 - Due by Mar. 21, 2022.