

Artificial Intelligence

Lecture 4: Constraint Satisfaction Problems

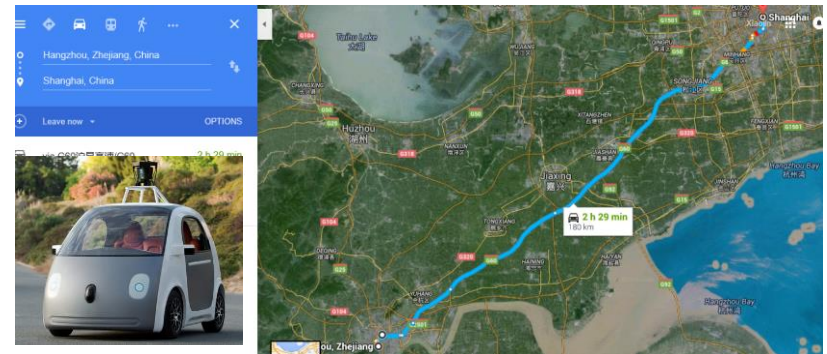

Xiaojin Gong

2022-03-14

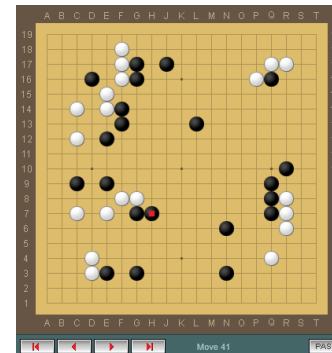
Credits: AI Courses in Berkeley

Review

- Search problems



- Adversarial search



Review

- Rational Agents
 - Structure
 - Representation

Search


Markov decision processes



(a) Atomic

Constraint satisfaction


Bayesian network



(b) Factored

Knowledge-based learning

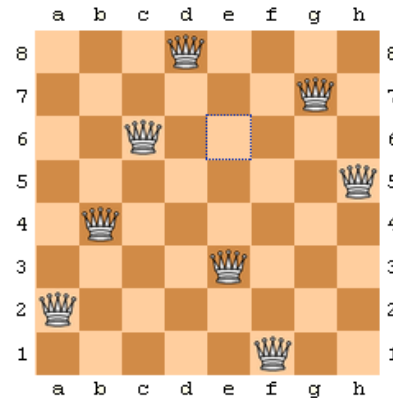
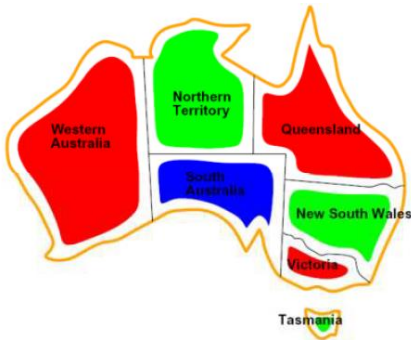
Natural language processing



(b) Structured

Outline

- Constraint satisfaction problems
 - Problem formulation
 - Backtracking search
 - Inference in CSPs



	5	3	2		7		8
6		1	5				2
2			9	1	3		5
7	1	4	6	9	2		
	2					6	
			4	5	1	2	9
	6		3	2	5		9
1					6	3	4
8			1	9	6	7	

- A special subset of search problems
- The goal itself is important, not the path

Constraint Satisfaction Problem

- A CPS consists of three components:
 - A set of **variables** $X = \{X_1, \dots, X_n\}$
 - A set of **domains** $D = \{D_1, \dots, D_n\}$
 - A set of **constraints** C
- **State** is defined by an **assignment** of values from a domain D_i to some or all of the variables X_i .
 - An assignment that does not violate any constraints is called a **consistent** or **legal assignment**.
 - A **complete assignment** is one in which every variable is assigned.
- **Goal test** is a set of constraints specifying allowable combinations of values for subsets of variables.
- **Solution** is a **consistent, complete assignment**.



Example: Map Coloring

- **Variables:** $X = \{WA, NT, SA, Q, NSW, V, T\}$
- **Domains:** $D = \{r, g, b\}$
- **Constraints:**
 - Implicit: $WA \neq NT, \dots$
 - Explicit: $(WA, NT) \in \{(r, g), (r, b), \dots\},$
...
- **Solutions** are assignments satisfying all constraints



Constraint Graph


- Each node corresponds to a variable
- Each arc participates in a constraint
- General-purpose CSP algorithms use the graph structure to speed up search.



Example: Cryptarithmic Puzzle

- Variables: $X = \{F, T, U, W, R, O, C_1, C_2, C_3\}$
- Domains: $D = \{0, 1, 2, \dots, 9\}$
- Constraints:
 - $\text{Alldiff}(F, T, U, W, R, O)$
 - $O + O = R + 10C_1$
 - $C_1 + W + W = U + 10C_2$
 - $C_2 + T + T = O + 10C_3$
 - $C_3 = F$

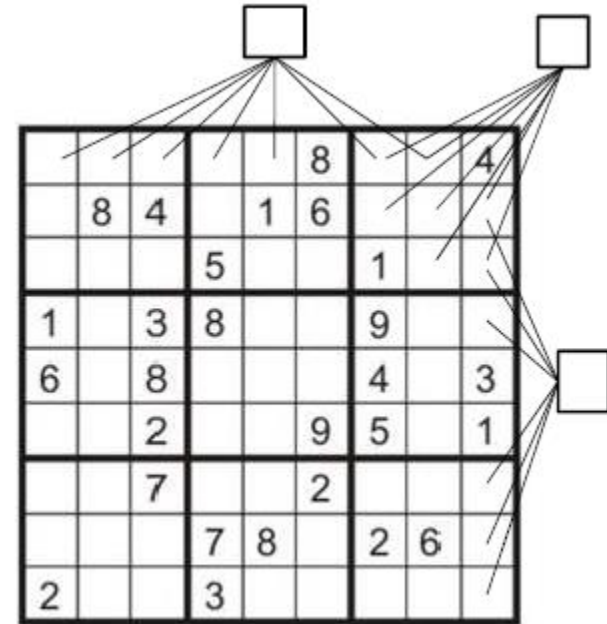
T	W	O
$+$	T	W
<hr/>		
F	O	U
		R



Constraint Hypergraph

Example: Sudoku

- **Variables:** Each (open) square
- **Domains:** $D = \{1, 2, \dots, 9\}$
- **Constraints:**
 - 9-way *Alldiff* for each column
 - 9-way *Alldiff* for each row
 - 9-way *Alldiff* for each region
- **Solutions** are assignments satisfying all constraints



Example: Job-shop Scheduling

- Car assembly tasks:
 - Install axles (front and back), affix all four wheels (right and left, front and back), tighten nuts for each wheel, affix hubcaps, and inspect the final assembly

- Variables:

$$X = \{Axle_F, Axle_B, Wheel_{RF}, Wheel_{LF}, Wheel_{RB}, Wheel_{LB}, Nuts_{RF}, Nuts_{LF}, Nuts_{RB}, Nuts_{LB}, Cap_{RF}, Cap_{LF}, Cap_{RB}, Cap_{LB}, Inspect\} .$$

- Domains: $D = \{1, 2, \dots, 27\}$

- Constraints:

- Precedence constraint

$$Axle_F + 10 \leq Wheel_{RF}; \quad Wheel_{RF} + 1 \leq Nuts_{RF}; \quad Nuts_{RF} + 2 \leq Cap_{RF};$$

- Disjunctive constraint

$$(Axle_F + 10 \leq Axle_B) \quad \text{or} \quad (Axle_B + 10 \leq Axle_F)$$

Example: Real World CSPs

- Timetabling problems
- Transportation scheduling
- Factory scheduling
- Hardware configuration
- ...

Many real-world problems involve real-valued variables

Varieties of CSPs: Variables

- Discrete variables

- Finite domains

- E.g., map coloring, Sudoku

- Infinite domains



- E.g., job scheduling
 - Linear constraints solvable, nonlinear undecidable

- Continuous variables

- E.g., start/end times for Hubble Telescope observations
 - Linear constraints solvable in polynomial time by **linear programming** methods

Varieties of CSPs: Constraints

- Varieties of constraints
 - Unary constraints
 - E.g., map coloring, $SA \neq g$
 - Binary constraints
 - E.g., map coloring, $SA \neq WA$
 - Higher-order constraints
 - E.g., cryptarithmic column constraints
- Preference constraints (Soft constraints)
 - Indicating which solutions are preferred
 - E.g., red is better than green
 - Often represented by a cost for each variable assignment
 - Gives constrained optimization problems




Solving CSPs

- Standard search formulation
 - **States** defined by the values assigned so far
 - **Initial state**: the empty assignment
 - **Successor function**: assign a value to an unassigned variable
 - **Goal test**: if the assignment is complete and satisfies all constraints
- Search tree
 - For n variables of domain size d ,
 - Top level: nd
 - 2nd level: $(n-1)d$
 - Leaves: $n!d^n$
- Depth-first search

Commutativity in CPS!

Backtracking Search for CSPs

- Backtracking search is the basic uninformed algorithm for solving CSPs
- Backtracking search is the depth-first search with two improvements:
 1. One variable at a time
 - Variable assignments are commutative, so **fix ordering**
 - Only need to consider assignments to a single variable at each step
 2. Backtrack when a variable has no legal values
 - Incremental goal test
- The number of leaves in the search tree is d^n



Backtracking Search for CSPs

```
function BACKTRACKING-SEARCH(csp) returns a solution, or failure
  return BACKTRACK({ }, csp)


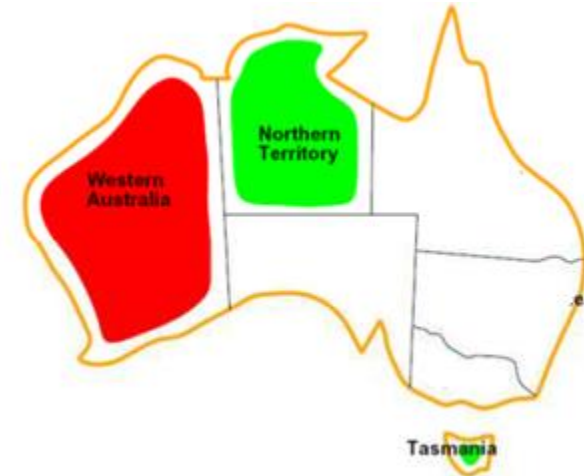
function BACKTRACK(assignment, csp) returns a solution, or failure
  if assignment is complete then return assignment
  var ← SELECT-UNASSIGNED-VARIABLE(csp)
  for each value in ORDER-DOMAIN-VALUES(var, assignment, csp) do
    if value is consistent with assignment then
      add { var = value } to assignment
      inferences ← INFERENCE(csp, var, value)
      if inferences ≠ failure then
        add inferences to assignment
        result ← BACKTRACK(assignment, csp)
        if result ≠ failure then
          return result
      remove { var = value } and inferences from assignment
  return failure
```


Backtracking Search for CSPs

- Ordering:
 - Which **variable** should be assigned next?
 - In what order its **values** be tried?
- Inference:
 - Infer reductions in the domain of variables
- Structure:
 - Can we exploit **the problem structure**?

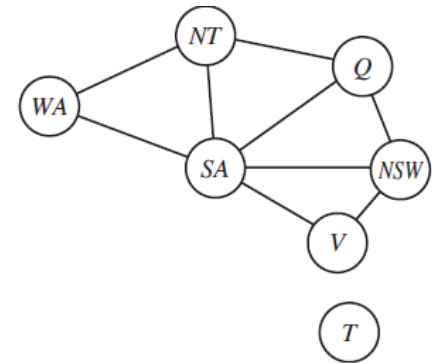
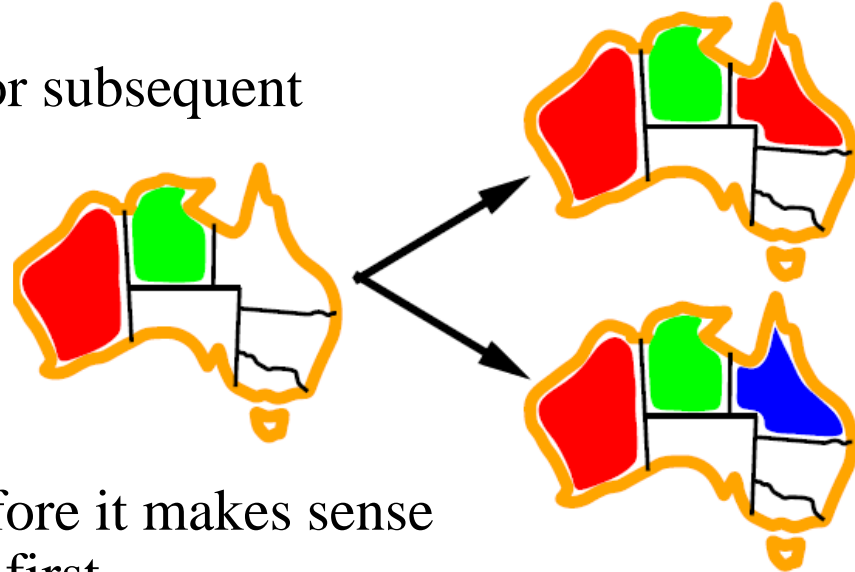
Variable and Value Ordering

- Variable ordering:
 - Minimum-remaining-values (MRV) heuristic
 - Choose the variable with the fewest “legal” values
 - Also called the **most constrained variable** or **fail-first heuristic**
 - Help minimize the number of nodes in the search tree by pruning larger parts of the tree earlier.
 - Degree heuristic
 - Choose the variable with the largest degree



Variable and Value Ordering

- Value ordering:
 - Least-constraining-value heuristic
 - Leave the maximum flexibility for subsequent variable assignments
 - Fail-last heuristic
 - We only need one solution, therefore it makes sense to look for the most likely values first.



Inference in CSPs

- Inference:
 - Constraint propagation
 - Reduce the number of legal values for a variable
 - Constraint propagation may be intertwined with search, or done as a preprocessing step
 - The key idea is **local consistency**
- Forms of inference:
 - Forward checking
 - Maintaining Arc Consistency (MAC)

```
for each value in ORDER-DOMAIN-VALUES(var, assignment, csp) do
  if value is consistent with assignment then
    add { var = value } to assignment
    inferences  $\leftarrow$  INFERENCE(csp, var, value)
```

Inference: Forward Checking

- Whenever a variable X is assigned, the forward-checking process establishes **arc consistency** for it





WA	NT	Q	NSW	V	SA
<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>
<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>
<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>

- Keep track of domains for unassigned variables and cross off bad options
- Forward checking propagates information from assigned to unassigned variables
- It makes the current variable arc-consistent, but doesn't look ahead and make all the other variables arc-consistent.

Arc Consistency

- An arc $X \rightarrow Y$ is consistent *iff* for **every** x there is **some** y which could be assigned without violating a constraint



Delete from the tail!

- Forward checking: Enforcing consistency of arcs pointing to each new assignment

Arc-Consistency Algorithm (AC-3)

function AC-3(*csp*) **returns** false if an inconsistency is found and true otherwise

inputs: *csp*, a binary CSP with components (X, D, C)

local variables: *queue*, a queue of arcs, initially all the arcs in *csp*

while *queue* is not empty **do**

$(X_i, X_j) \leftarrow \text{REMOVE-FIRST}(\text{queue})$


if REVISE(*csp*, X_i , X_j) **then**

if size of $D_i = 0$ **then return** false

for each X_k **in** $X_i.\text{NEIGHBORS} - \{X_j\}$ **do**

 add (X_k, X_i) to *queue*

return true



function REVISE(*csp*, X_i , X_j) **returns** true iff we revise the domain of X_i

revised \leftarrow false

for each x **in** D_i **do**

if no value y in D_j allows (x, y) to satisfy the constraint between X_i and X_j **then**

 delete x from D_i

revised \leftarrow true


return *revised*

Arc-Consistency Algorithm (AC-3)

- Runtime:
 - Assume a CSP with n variables, d domain size, and c binary constraints
 - Each arc (X_k, X_i) can be inserted d times
 - Checking consistency of an arc takes $O(d^2)$ time
 - So the total worst-case time is $O(cd^3)$

Arc-Consistency Algorithm (AC-3)



- A simple form of propagation makes sure all arcs are consistent:



- Important: If X loses a value, neighbors of X need to be rechecked!
- Arc consistency detects failure earlier than forward checking
- Can be run as a preprocessor or after each assignment
- Interleaving inference with search

Limitations of Arc Consistency

- After enforcing arc consistency:
 - Can have one solution left
 - Can have multiple solutions left
 - Can have no solutions left (and not know it)
- Arc consistency still runs inside a backtracking search!



K-Consistency


- Increasing degrees of consistency
 - 1-Consistency (Node Consistency):
 - Each single node's domain has a value which meets that node's unary constraints
 - 2-Consistency (Arc Consistency):
 - For each pair of nodes, any consistent assignment to one can be extended to the other
 - K-Consistency:
 - For each k nodes, any consistent assignment to $k-1$ can be extended to the k th node.
- Higher k more expensive to compute

Strong K-Consistency

- Strong k-consistency:
 - Also k-1, k-2, ... 1 consistent
- Strong n-consistency means we can solve without backtracking
 - Guaranteed to find a solution in $O(n^2d)$
 - Choose any assignment to any variable
 - Choose a new variable
 - By 2-consistency, there is a choice consistent with the first
 - Choose a new variable
 - By 3-consistency, there is a choice consistent with the first 2
 - ...


The Structure of Problems

- Decompose into independent subproblems
 - Independence can be ascertained by finding connected components
- Suppose each subproblem CSP_i has c variables from the total of n variables,
 - The total work is $O(n/c d^c)$
 - E.g. $n=80, d=2, c=20$
 - $2^{80} = 4$ billion years at 10 million nodes/sec
 - $4 \cdot 2^{20} = 0.4$ seconds at 10 million nodes/sec.



Tree-Structured CSP

- Any two variables are connected by only one path
- Any tree-structured CSP can be solved in $O(nd^2)$ time
 - Compare to general CSPs, where worst-case time is $O(d^n)$



Tree-Structured CSP

function TREE-CSP-SOLVER(*csp*) **returns** a solution, or failure

inputs: *csp*, a CSP with components X , D , C

$n \leftarrow$ number of variables in X

assignment \leftarrow an empty assignment

root \leftarrow any variable in X

$X \leftarrow$ TOPOLOGICALSORT(X , *root*)

for $j = n$ **down to** 2 **do**

MAKE-ARC-CONSISTENT(PARENT(X_j), X_j)

if it cannot be made consistent **then return** *failure*

for $i = 1$ **to** n **do**

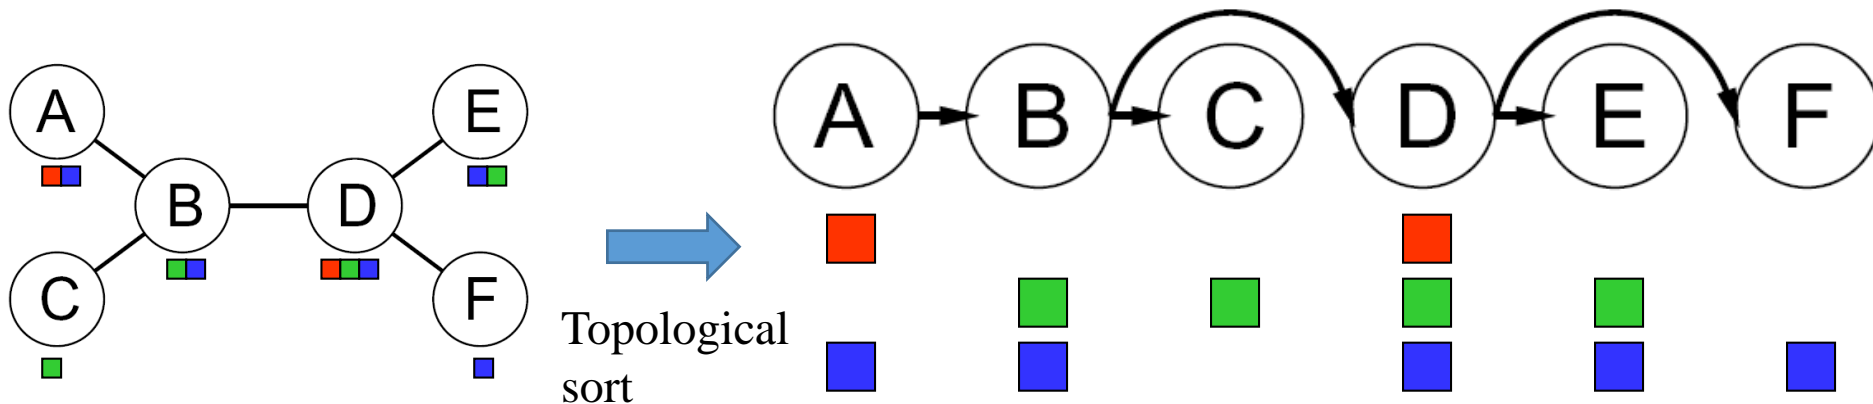
assignment[X_i] \leftarrow any consistent value from D_i

if there is no consistent value **then return** *failure*

return *assignment*

Tree-Structured CSP

- Topological sort
 - Each variable appears after its parent in the tree.
- Directed arc consistency (DAC)
 - A CSP is DAC under an ordering X_1, X_2, \dots, X_n iff every X_i is arc-consistent with each X_j for $j > i$




- No backtrack!

Nearly Tree-Structured CSP

- Reduce general constraint graphs to trees
 1. Removing nodes
 - **Cutset conditioning**: instantiate (in all ways) a set of variables such that the remaining constraint graph is a tree

1. Choose a subset S of the CSP's variables such that the constraint graph becomes a tree after removal of S . S is called a **cycle cutset**.
2. For each possible assignment to the variables in S that satisfies all constraints on S ,
 - (a) remove from the domains of the remaining variables any values that are inconsistent with the assignment for S , and
 - (b) If the remaining CSP has a solution, return it together with the assignment for S .



Nearly Tree-Structured CSP


- Reduce general constraint graphs to trees
 1. Removing nodes
 - **Cutset conditioning**: instantiate (in all ways) a set of variables such that the remaining constraint graph is a tree
 - Cutset size c gives runtime $O(d^c (n-c) d^2)$, very fast for small c

1. Choose a subset S of the CSP's variables such that the constraint graph becomes a tree after removal of S . S is called a **cycle cutset**.
2. For each possible assignment to the variables in S that satisfies all constraints on S ,
 - (a) remove from the domains of the remaining variables any values that are inconsistent with the assignment for S , and
 - (b) If the remaining CSP has a solution, return it together with the assignment for S .

Nearly Tree-Structured CSP

- Reduce general constraint graphs to trees
 2. Collapsing nodes together
 - **Tree decomposition:** construct the constraint graph into a set of connected subproblems.

- Every variable in the original problem appears in at least one of the subproblems.
- If two variables are connected by a constraint in the original problem, they must appear together (along with the constraint) in at least one of the subproblems.
- If a variable appears in two subproblems in the tree, it must appear in every subproblem along the path connecting those subproblems.



Nearly Tree-Structured CSP

- Reduce general constraint graphs to trees
 - 2. Collapsing nodes together
 - View each subproblem as a “mega-variable” whose domain is the set of all solutions for the subproblem.
 - Solve the constraints connecting between subproblems using the algorithm for tree-structured CSPs.
 - The problem can be solved in $O(nd^{w+1})$, w is the tree width of a tree decomposition.

Assignments

- Reading assignment:
 - Ch. 6.1-6.3, 6.5